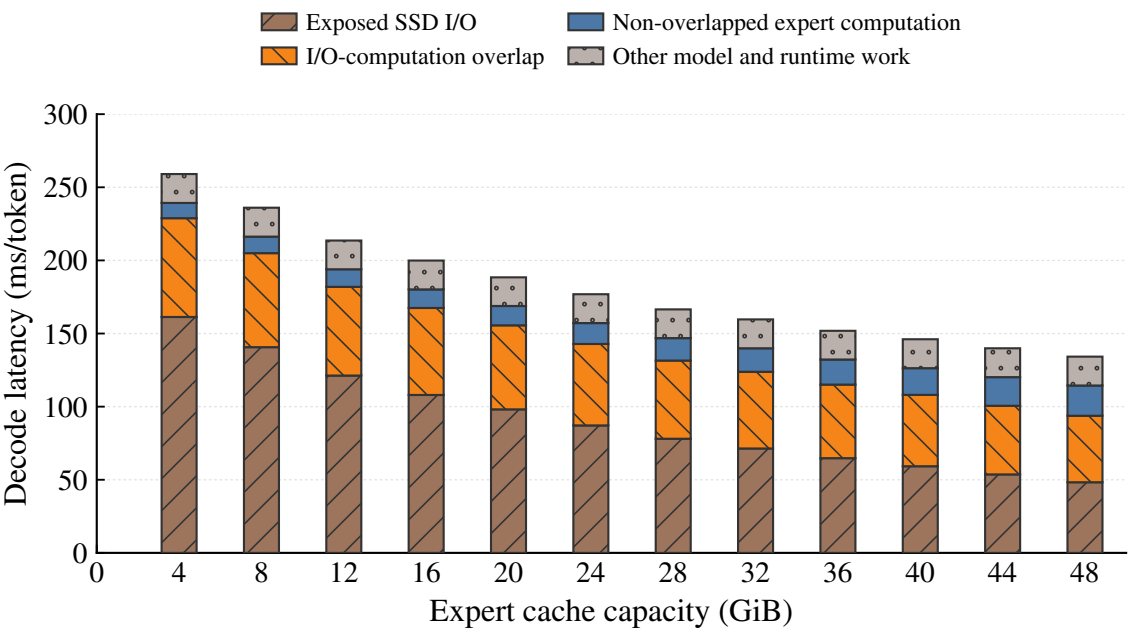
Exposed SSD I/O
I/O-computation overlap
Non-overlapped expert computation
Other model and runtime work
Decode latency (ms/token)
300
250
200
150
100
50
0
0
4
8
12
16
20
24
28
32
36
40
44
48
Expert cache capacity (GiB)