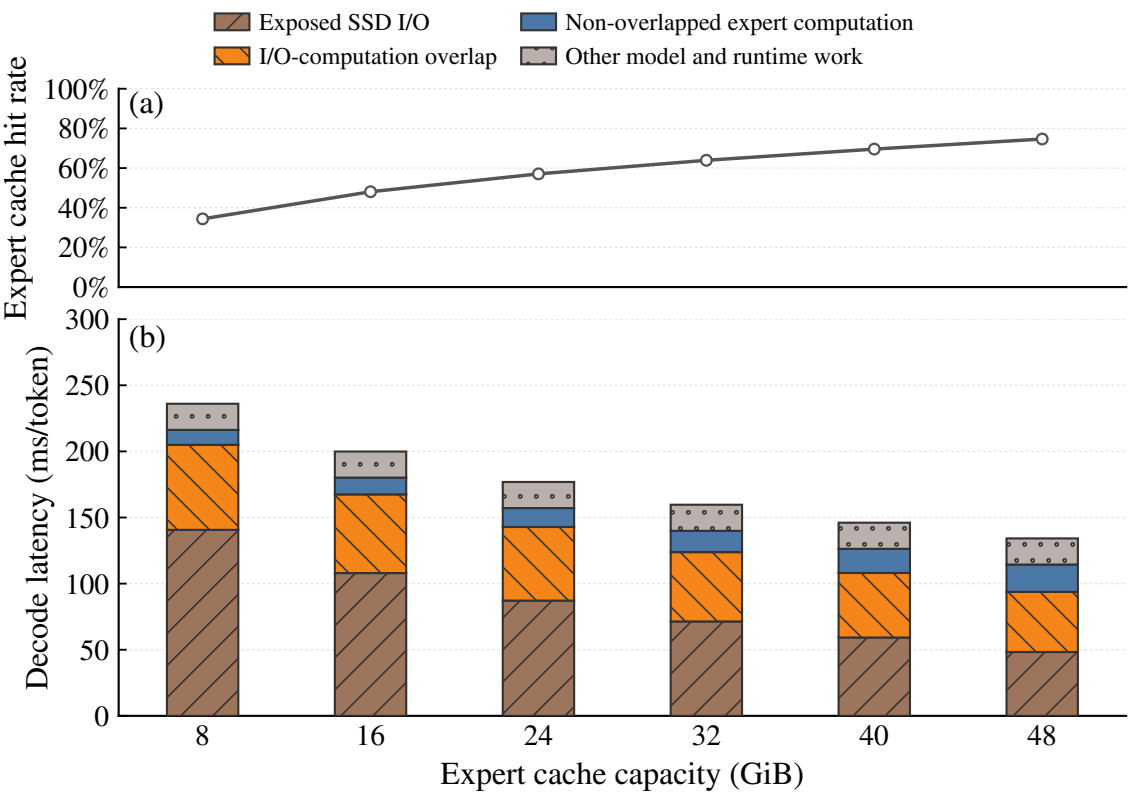
Exposed SSD I/O
Non-overlapped expert computation
I/O-computation overlap
Other model and runtime work
(a)
Expert cache hit rate
100%
80%
60%
40%
20%
0%
(b)
Decode latency (ms/token)
300
250
200
150
100
50
0
8
16
24
32
40
48
Expert cache capacity (GiB)