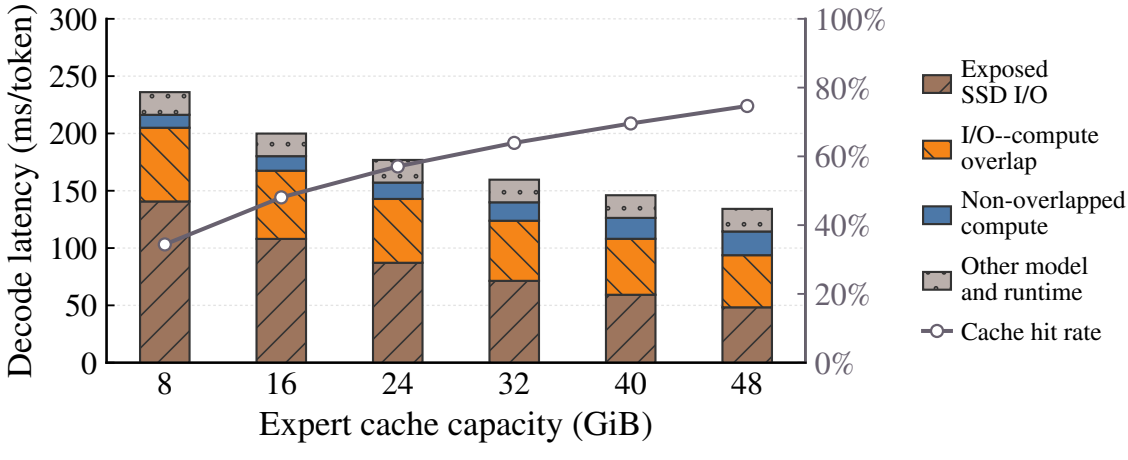
Decode latency (ms/token)
300
250
200
150
100
50
0
100%
80%
60%
40%
20%
0%
8
16
24
32
40
48
Expert cache capacity (GiB)
Exposed SSD I/O
I/O--compute overlap
Non-overlapped compute
Other model and runtime
Cache hit rate