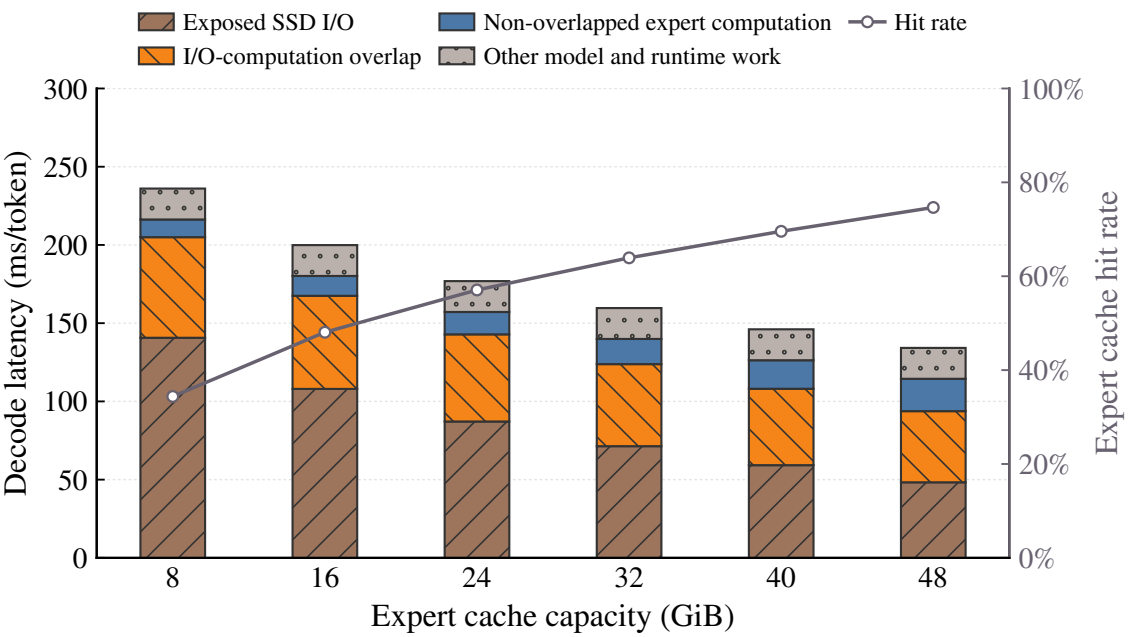
Exposed SSD I/O
I/O-computation overlap
Non-overlapped expert computation
Other model and runtime work
Hit rate
Decode latency (ms/token)
Expert cache hit rate
Expert cache capacity (GiB)
0
50
100
150
200
250
300
0%
20%
40%
60%
80%
100%
8
16
24
32
40
48